

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv("https://raw.githubusercontent.com/ektanegi25/Water-Sensor-repositor
```

```
In [79]: df.head(5)
```

```
Out[79]:
```

	Wafers	Sensor-1	Sensor-2	Sensor-3	Sensor-4	Sensor-5	Sensor-6	Sensor-7	Sensor-8	Sensor-9
55	Wafer-856	NaN	2532.45	2191.1333	2197.6570	1.1569	100.0	89.7222	0.1251	1.5762
88	Wafer-889	3221.21	2391.20	2189.9667	1046.6212	0.8662	100.0	102.3622	0.1208	1.4756
26	Wafer-827	2951.85	2525.00	2189.5777	1320.3197	1.3459	100.0	100.7744	0.1234	1.5590
42	Wafer-843	2982.07	2447.06	2199.6334	1242.8420	1.4083	100.0	99.2178	0.1221	1.4542
69	Wafer-870	3058.08	2524.60	2192.3778	1110.5453	0.8147	100.0	99.2922	0.1226	1.4958

5 rows × 592 columns

```
In [4]: df.tail(5)
```

```
Out[4]:
```

	Unnamed: 0	Sensor-1	Sensor-2	Sensor-3	Sensor-4	Sensor-5	Sensor-6	Sensor-7	Sensor-8	Sensor-9
95	Wafer-896	3013.66	2526.44	2185.2111	1141.6306	0.8447	100.0	100.5978	0.1217	1.533
96	Wafer-897	2982.87	2477.01	2315.2667	2360.1325	1.1259	100.0	90.1144	0.1160	1.465
97	Wafer-898	3084.82	2387.42	2171.5000	1028.4440	0.7899	100.0	101.5122	0.1224	1.360
98	Wafer-899	2955.87	2541.89	NaN	NaN	NaN	NaN	NaN	NaN	1.445
99	Wafer-900	2914.86	2465.11	2210.2778	2120.5760	1.0700	100.0	95.1089	0.1230	1.581

5 rows × 592 columns

```
In [5]: df.rename(columns = {'Unnamed: 0':"Wafers"}, inplace = True)
```

```
In [6]: df.head(5)
```

Out[6]:

	Wafers	Sensor-1	Sensor-2	Sensor-3	Sensor-4	Sensor-5	Sensor-6	Sensor-7	Sensor-8	Sensor-9	..
0	Wafer-801	2968.33	2476.58	2216.7333	1748.0885	1.1127	100.0	97.5822	0.1242	1.5300	.
1	Wafer-802	2961.04	2506.43	2170.0666	1364.5157	1.5447	100.0	96.7700	0.1230	1.3953	.
2	Wafer-803	3072.03	2500.68	2205.7445	1363.1048	1.0518	100.0	101.8644	0.1220	1.3896	.
3	Wafer-804	3021.83	2419.83	2205.7445	1363.1048	1.0518	100.0	101.8644	0.1220	1.4108	.
4	Wafer-805	3006.95	2435.34	2189.8111	1084.6502	1.1993	100.0	104.8856	0.1234	1.5094	.

5 rows × 592 columns

```
In [7]: from sklearn.model_selection import train_test_split
```

```
In [8]: df,df_test=train_test_split(df,test_size=0.2,random_state=42)
```

```
In [9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 80 entries, 55 to 51
Columns: 592 entries, Wafers to Good/Bad
dtypes: float64(494), int64(97), object(1)
memory usage: 370.6+ KB
```

```
In [10]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20 entries, 83 to 31
Columns: 592 entries, Wafers to Good/Bad
dtypes: float64(494), int64(97), object(1)
memory usage: 92.7+ KB
```

```
In [11]: df.describe()
```

Out[11]:

	Sensor-1	Sensor-2	Sensor-3	Sensor-4	Sensor-5	Sensor-6	Sensor-7	Sensor-8
count	79.000000	80.000000	78.000000	78.000000	78.000000	78.0	78.000000	78.000000
mean	3019.048228	2494.058875	2202.758988	1519.467071	1.201382	100.0	96.881160	0.122000
std	72.665372	68.166898	31.633772	471.962104	0.365870	0.0	5.520108	0.002000
min	2889.670000	2254.990000	2114.666700	978.783200	0.753100	100.0	83.423300	0.116000
25%	2975.425000	2452.517500	2189.966700	1111.543600	0.850075	100.0	93.547250	0.120000
50%	3004.390000	2502.445000	2200.955600	1308.647900	1.164250	100.0	99.217800	0.122000
75%	3065.730000	2532.755000	2212.866700	1997.641600	1.383000	100.0	101.133300	0.123000
max	3221.210000	2664.520000	2315.266700	2363.641200	2.207300	100.0	103.091100	0.126000

8 rows × 591 columns

In [12]: df['Good/Bad'].isnull().sum()

Out[12]:

0

In [13]: df.isnull().sum()

Out[13]:

Wafers	0
Sensor-1	1
Sensor-2	0
Sensor-3	2
Sensor-4	2
..	
Sensor-587	0
Sensor-588	0
Sensor-589	0
Sensor-590	0
Good/Bad	0
Length: 592, dtype: int64	

In [14]: df.isnull().sum().sum()

Out[14]:

1822

In [15]: df.isnull().sum().sum() / (df.shape[0] * (df.shape[1] - 1)) * 100

Out[15]:

3.8536379018612523

In [16]: df_test.isnull().sum().sum()

Out[16]:

484

In [17]: df.head()

Out[17]:

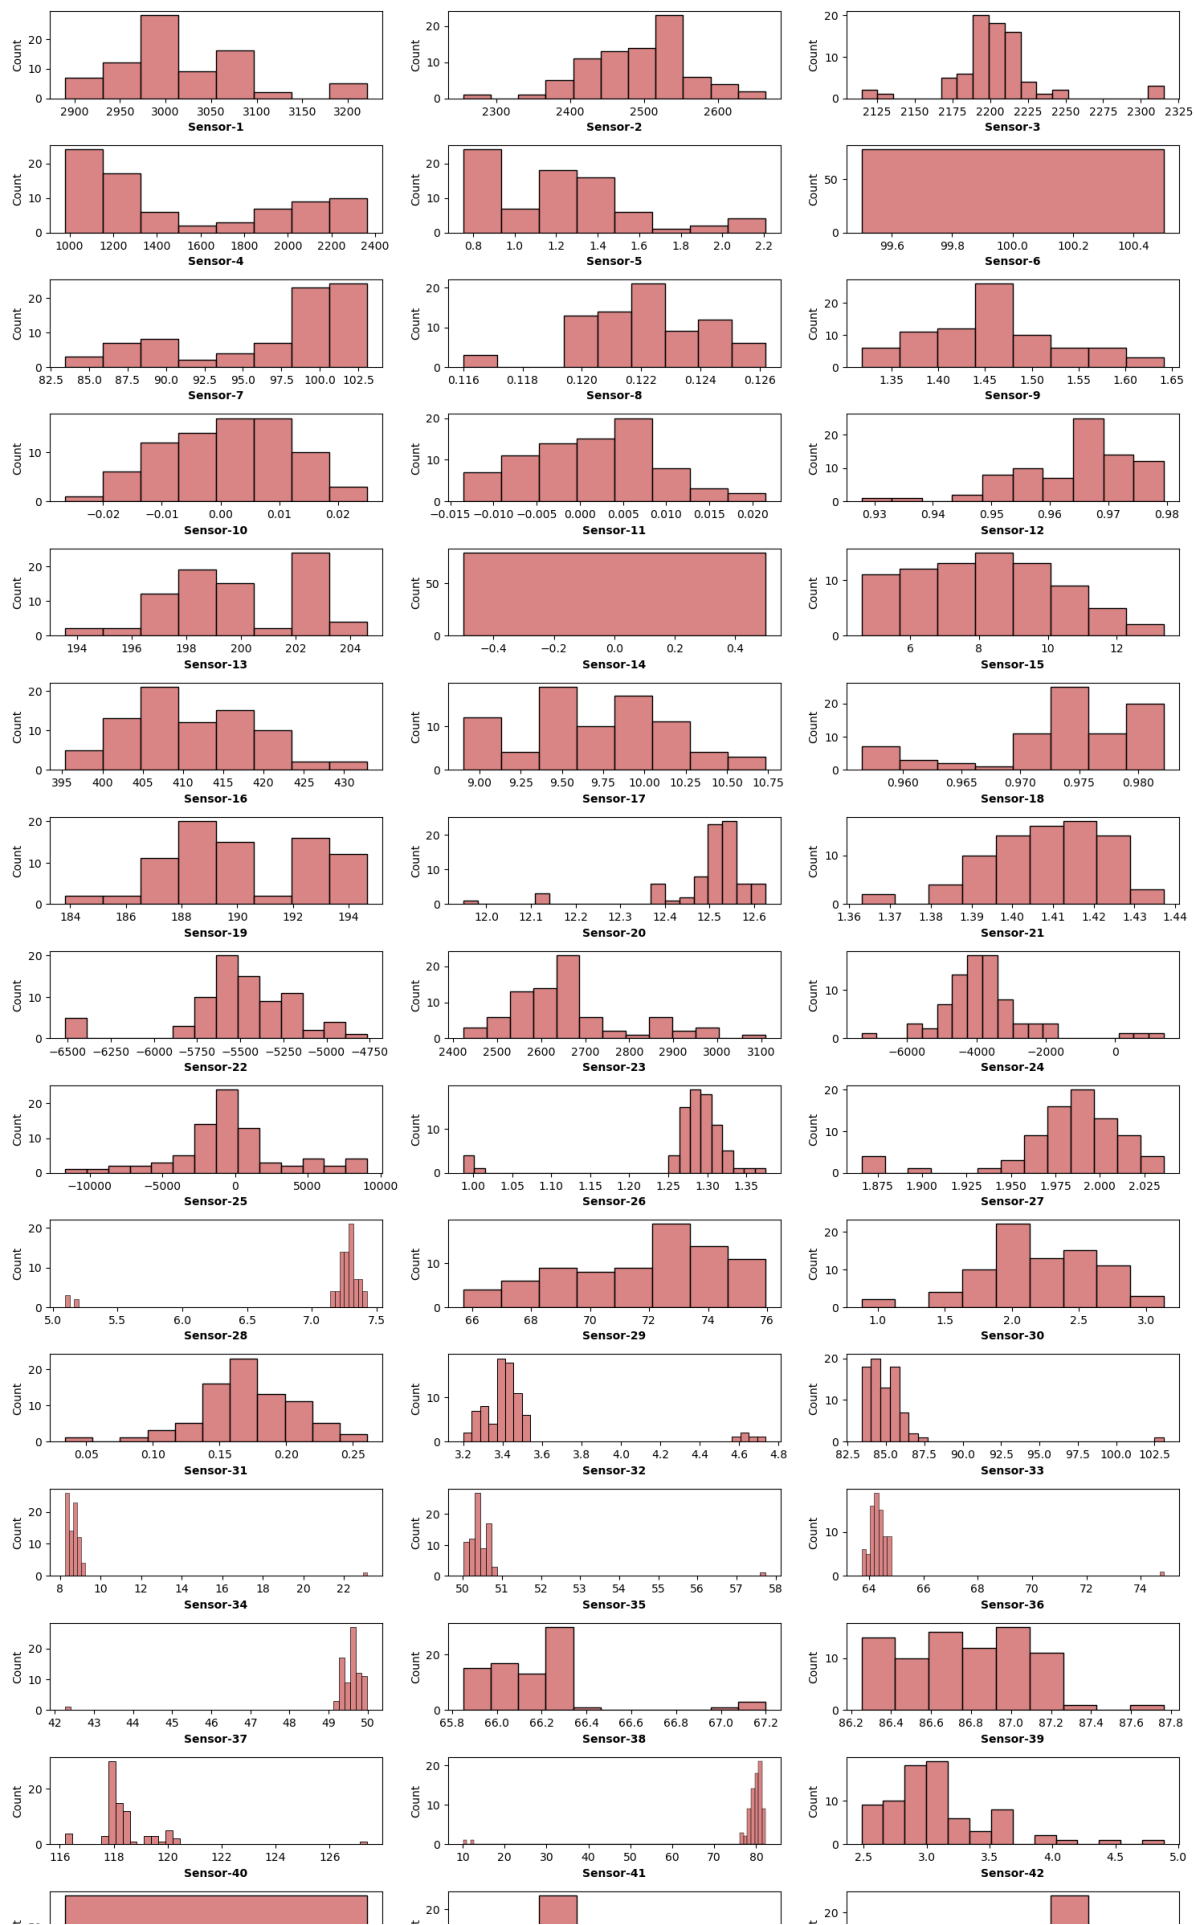
	Wafers	Sensor-1	Sensor-2	Sensor-3	Sensor-4	Sensor-5	Sensor-6	Sensor-7	Sensor-8	Sensor-9
55	Wafer-856	NaN	2532.45	2191.1333	2197.6570	1.1569	100.0	89.7222	0.1251	1.5762
88	Wafer-889	3221.21	2391.20	2189.9667	1046.6212	0.8662	100.0	102.3622	0.1208	1.4756
26	Wafer-827	2951.85	2525.00	2189.5777	1320.3197	1.3459	100.0	100.7744	0.1234	1.5590
42	Wafer-843	2982.07	2447.06	2199.6334	1242.8420	1.4083	100.0	99.2178	0.1221	1.4542
69	Wafer-870	3058.08	2524.60	2192.3778	1110.5453	0.8147	100.0	99.2922	0.1226	1.4958

5 rows × 592 columns

```

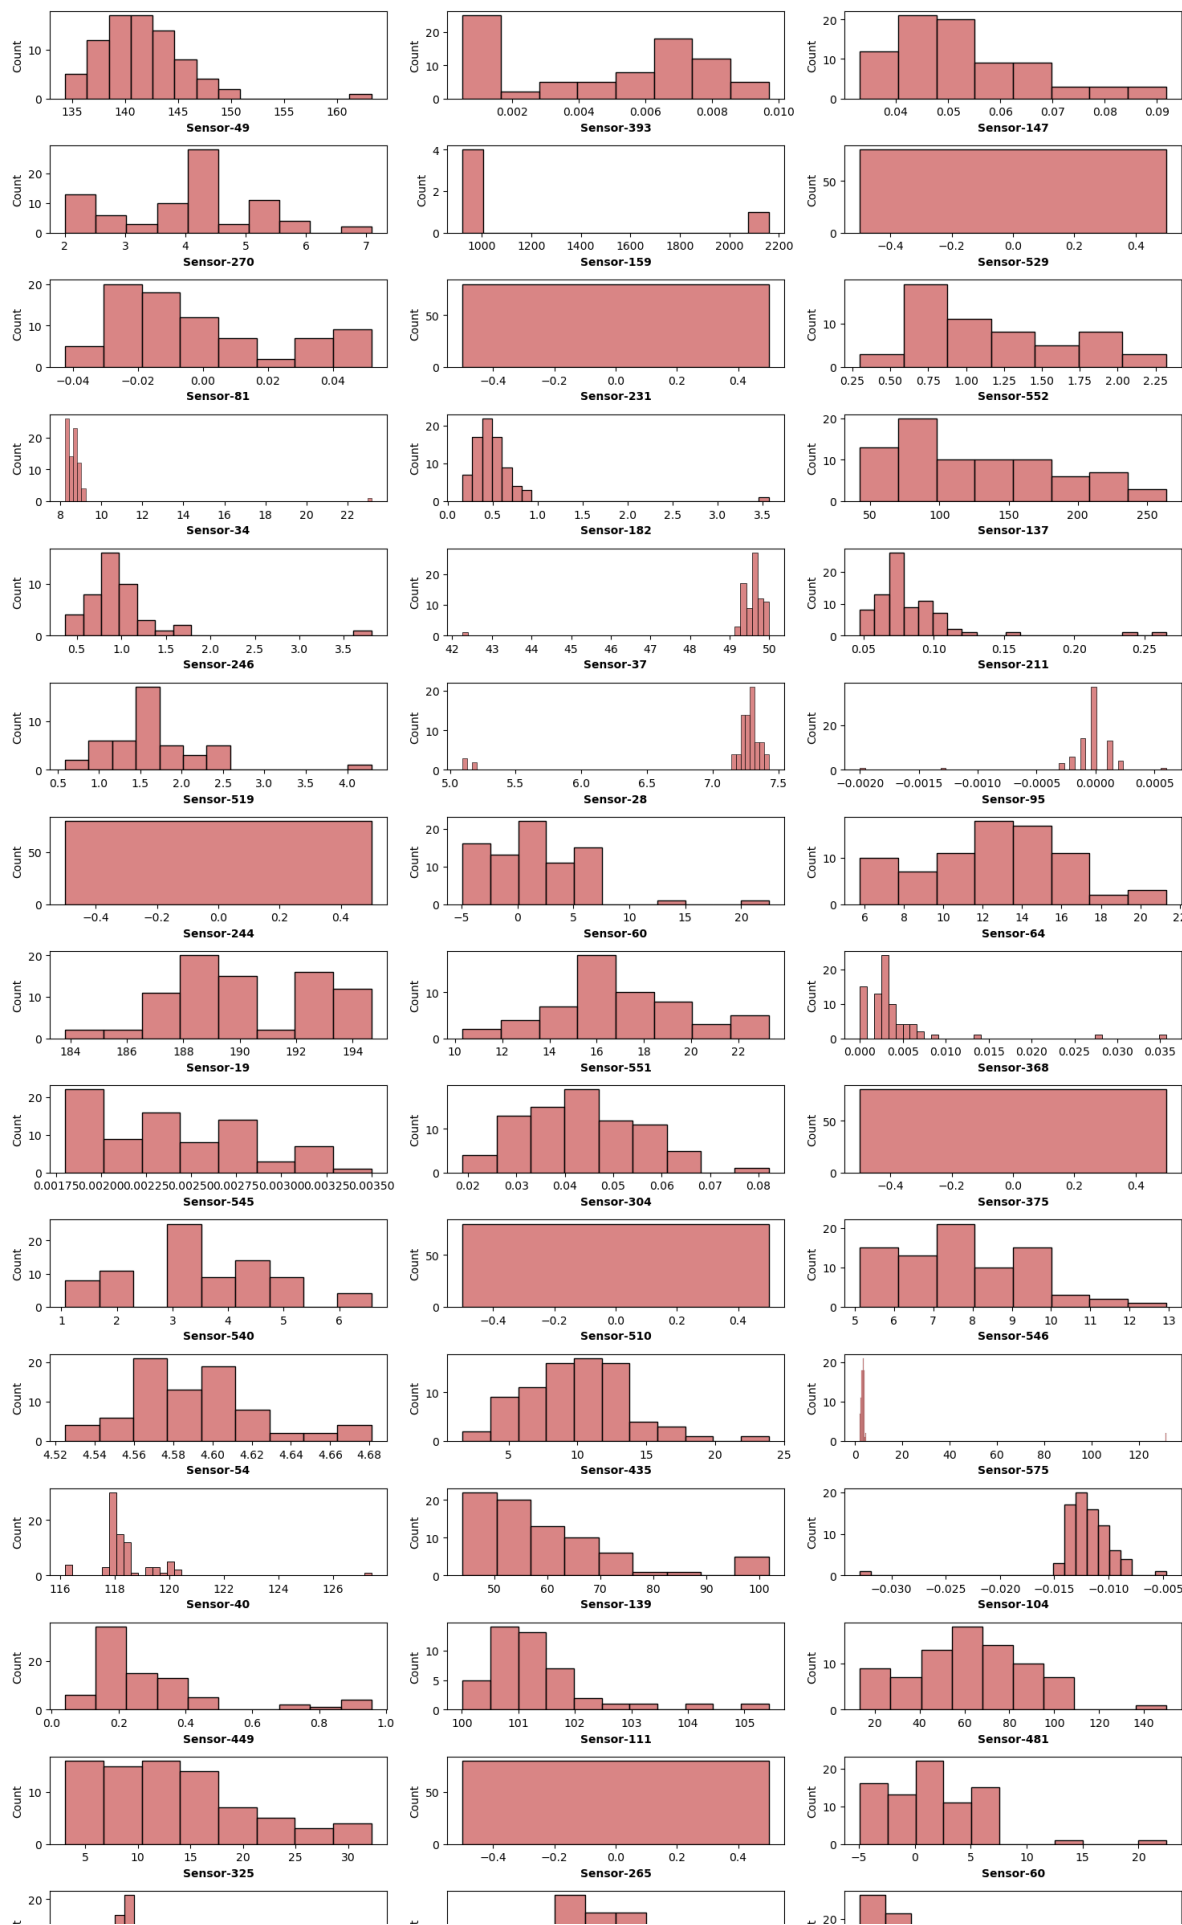
In [18]: plt.figure(figsize=(15,100))
for i,col in enumerate(df.columns[1:51]):
    plt.subplot(60,3,i+1)
    sns.histplot(x=df[col],color='indianred')
    plt.xlabel(col,weight='bold')
    plt.tight_layout()

```



```
In [22]: random_50_sensors=[]  
for i in range(50):  
    if i not in random_50_sensors:  
        random_50_sensors.append(np.random.randint(1,591))
```

```
In [20]: #lets have a look to first 50 sensors  
  
plt.figure(figsize=(15,100))  
for i,col in enumerate(df.columns[random_50_sensors]):  
    plt.subplot(60,3,i+1)  
    sns.histplot(x=df[col],color='indianred')  
    plt.xlabel(col,weight='bold')  
    plt.tight_layout()
```



```
In [23]: def get_cols_zero_std(df:pd.DataFrame):  
        cols_to_drop=[]  
        num_cols=[i for i in df.columns[1:] if df[i].dtype != "0"]  
        for i in df.columns[1:]:  
            if df[i].std()==0:  
                cols_to_drop.append(i)  
        return cols_to_drop
```

```
In [24]: def get_redundant_col(df:pd.DataFrame,missing_thresh=.7):  
        cols_missing_ratio=df.isnull().sum().div(df.shape[0])  
        cols_to_drop = list(cols_missing_ratio[cols_missing_ratio > missing_thresh].ind  
        return cols_to_drop
```

```
In [25]: cols_drop_1=get_cols_zero_std(df=df)
```

```
In [26]: cols_drop_1
```



```
Out[26]: ['Sensor-6',  
          'Sensor-14',  
          'Sensor-43',  
          'Sensor-50',  
          'Sensor-53',  
          'Sensor-70',  
          'Sensor-75',  
          'Sensor-98',  
          'Sensor-142',  
          'Sensor-150',  
          'Sensor-179',  
          'Sensor-180',  
          'Sensor-187',  
          'Sensor-190',  
          'Sensor-191',  
          'Sensor-192',  
          'Sensor-193',  
          'Sensor-194',  
          'Sensor-195',  
          'Sensor-207',  
          'Sensor-210',  
          'Sensor-227',  
          'Sensor-230',  
          'Sensor-231',  
          'Sensor-232',  
          'Sensor-233',  
          'Sensor-234',  
          'Sensor-235',  
          'Sensor-236',  
          'Sensor-237',  
          'Sensor-238',  
          'Sensor-241',  
          'Sensor-242',  
          'Sensor-243',  
          'Sensor-244',  
          'Sensor-257',  
          'Sensor-258',  
          'Sensor-259',  
          'Sensor-260',  
          'Sensor-261',  
          'Sensor-262',  
          'Sensor-263',  
          'Sensor-264',  
          'Sensor-265',  
          'Sensor-266',  
          'Sensor-267',  
          'Sensor-277',  
          'Sensor-285',  
          'Sensor-314',  
          'Sensor-315',  
          'Sensor-316',  
          'Sensor-323',  
          'Sensor-326',  
          'Sensor-327',  
          'Sensor-328',  
          'Sensor-329',
```

'Sensor-330',
'Sensor-331',
'Sensor-343',
'Sensor-348',
'Sensor-365',
'Sensor-370',
'Sensor-371',
'Sensor-372',
'Sensor-373',
'Sensor-374',
'Sensor-375',
'Sensor-376',
'Sensor-379',
'Sensor-380',
'Sensor-381',
'Sensor-382',
'Sensor-395',
'Sensor-396',
'Sensor-397',
'Sensor-398',
'Sensor-399',
'Sensor-400',
'Sensor-401',
'Sensor-402',
'Sensor-403',
'Sensor-404',
'Sensor-405',
'Sensor-415',
'Sensor-423',
'Sensor-450',
'Sensor-451',
'Sensor-452',
'Sensor-459',
'Sensor-462',
'Sensor-463',
'Sensor-464',
'Sensor-465',
'Sensor-466',
'Sensor-467',
'Sensor-479',
'Sensor-482',
'Sensor-499',
'Sensor-502',
'Sensor-503',
'Sensor-504',
'Sensor-505',
'Sensor-506',
'Sensor-507',
'Sensor-508',
'Sensor-509',
'Sensor-510',
'Sensor-513',
'Sensor-514',
'Sensor-515',
'Sensor-516',
'Sensor-529',

```
'Sensor-530',  
'Sensor-531',  
'Sensor-532',  
'Sensor-533',  
'Sensor-534',  
'Sensor-535',  
'Sensor-536',  
'Sensor-537',  
'Sensor-538',  
'Sensor-539']
```

```
In [27]: cols_drop_2=get_reduntant_col(df,missing_thresh=.7)
```

```
In [28]: cols_drop_2
```

```
Out[28]: ['Sensor-158', 'Sensor-159', 'Sensor-293', 'Sensor-294']
```

```
In [29]: cols_to_drop=cols_drop_1+cols_drop_2+['Wafers']
```

```
In [30]: len(cols_to_drop)
```

```
Out[30]: 127
```

```
In [31]: df
```

Out[31]:

	Wafers	Sensor-1	Sensor-2	Sensor-3	Sensor-4	Sensor-5	Sensor-6	Sensor-7	Sensor-8	Sensor-9
55	Wafer-856	NaN	2532.45	2191.1333	2197.6570	1.1569	100.0	89.7222	0.1251	1.5762
88	Wafer-889	3221.21	2391.20	2189.9667	1046.6212	0.8662	100.0	102.3622	0.1208	1.4756
26	Wafer-827	2951.85	2525.00	2189.5777	1320.3197	1.3459	100.0	100.7744	0.1234	1.5590
42	Wafer-843	2982.07	2447.06	2199.6334	1242.8420	1.4083	100.0	99.2178	0.1221	1.4542
69	Wafer-870	3058.08	2524.60	2192.3778	1110.5453	0.8147	100.0	99.2922	0.1226	1.4958
...
60	Wafer-861	3071.05	2642.15	2200.9889	1054.5240	1.3830	100.0	100.1800	0.1201	1.4532
71	Wafer-872	3043.18	2545.53	2192.3778	1110.5453	0.8147	100.0	99.2922	0.1226	1.3824
14	Wafer-815	3001.26	2519.92	2224.6778	1308.6479	1.3907	100.0	101.1333	0.1208	1.5172
92	Wafer-893	3007.00	2572.62	2213.2111	2070.7147	1.9705	100.0	87.7411	0.1232	1.4446
51	Wafer-852	3078.77	2533.04	2187.4111	1942.3069	1.1864	100.0	88.0911	0.1245	1.4500

80 rows × 592 columns



```
In [32]: X, y = df.drop(cols_to_drop, axis = 1), df['Good/Bad']
```

```
In [33]: X.shape
```

```
Out[33]: (80, 465)
```

```
In [34]: y.shape
```

```
Out[34]: (80,)
```

```
In [35]: y.value_counts()
```

```
Out[35]: -1    74
         1     6
         Name: Good/Bad, dtype: int64
```

```
In [36]: from sklearn.pipeline import Pipeline
         from sklearn.impute import KNNImputer, SimpleImputer
         from sklearn.preprocessing import RobustScaler
```

```
In [37]: imputer=KNNImputer(n_neighbors=3)
prep_process=Pipeline(steps=[('Imputer',imputer),('Scaling',RobustScaler())])
```

```
In [38]: prep_process
```

```
Out[38]: Pipeline
  ▸ KNNImputer
  ▸ RobustScaler
```

```
In [39]: X_trans = prep_process.fit_transform(X)
X_trans.shape
```

```
Out[39]: (80, 465)
```

```
In [40]: %pip install kneed
```

```
Collecting kneed
  Downloading kneed-0.8.5-py3-none-any.whl (10 kB)
Requirement already satisfied: scipy>=1.0.0 in /opt/conda/lib/python3.10/site-pack
ages (from kneed) (1.9.3)
Requirement already satisfied: numpy>=1.14.2 in /opt/conda/lib/python3.10/site-pac
kages (from kneed) (1.23.5)
Installing collected packages: kneed
Successfully installed kneed-0.8.5
Note: you may need to restart the kernel to use updated packages.
```

```
In [41]: from sklearn.cluster import KMeans
from kneed import KneeLocator
import numpy as np
from typing import Tuple

def cluster_data_instances(X: np.array) -> Tuple[KMeans, np.array]:
    wcss = [] # Within Summation of Squares

    for i in range(1, 11):
        kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
        kmeans.fit(X)
        wcss.append(kmeans.inertia_)

    knee_finder = KneeLocator(
        range(1, 11), wcss, curve='convex', direction='decreasing')
    ideal_clusters = knee_finder.knee

    kmeans = KMeans(n_clusters=ideal_clusters, init='k-means++', random_state=42)
    y_kmeans = kmeans.fit_predict(X)

    return kmeans, np.c_[X, y_kmeans]

# Usage example:
kmeans, X_clus = cluster_data_instances(X_trans)
```

```

/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(
/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870: FutureWarn
ing: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the v
alue of `n_init` explicitly to suppress the warning
  warnings.warn(

```

In [42]: kmeans

Out[42]: **KMeans**
 KMeans(n_clusters=3, random_state=42)

In [43]: X_clus

```
Out[43]: array([[ -0.02781221,  0.37395233, -0.4289214 , ...,  0.75955556,
                0.          , 1.          ],
               [ 2.50431022, -1.38644649, -0.47986463, ..., -0.52610857,
                0.          , 1.          ],
               [-0.60204699,  0.28110298, -0.49685153, ..., -0.08775867,
                0.          , 1.          ],
               ...,
               [-0.03223295,  0.21779093,  1.03590393, ..., -0.17805529,
                0.          , 2.          ],
               [ 0.03396281,  0.87459106,  0.53517467, ...,  0.28421459,
                0.          , 1.          ],
               [ 0.86164048,  0.3813055 , -0.59146288, ..., -0.20782888,
                0.          , 1.          ]])
```

```
In [44]: np.unique(X_clus[:, -1])
```

```
Out[44]: array([0., 1., 2.])
```

```
In [46]: wafer_clus = np.c_[X_clus, y]
        wafer_clus
```

```
Out[46]: array([[ -0.02781221,  0.37395233, -0.4289214 , ...,  0.          ,
                1.          , -1.          ],
               [ 2.50431022, -1.38644649, -0.47986463, ...,  0.          ,
                1.          , -1.          ],
               [-0.60204699,  0.28110298, -0.49685153, ...,  0.          ,
                1.          , -1.          ],
               ...,
               [-0.03223295,  0.21779093,  1.03590393, ...,  0.          ,
                2.          , -1.          ],
               [ 0.03396281,  0.87459106,  0.53517467, ...,  0.          ,
                1.          , -1.          ],
               [ 0.86164048,  0.3813055 , -0.59146288, ...,  0.          ,
                1.          , -1.          ]])
```

```
In [47]: wafer_clus[wafer_clus[:, -2] == 0].shape
```

```
Out[47]: (1, 467)
```

```
In [48]: wafer_clus[wafer_clus[:, -2] == 1].shape
```

```
Out[48]: (62, 467)
```

```
In [49]: wafer_clus[wafer_clus[:, -2] == 2].shape
```

```
Out[49]: (17, 467)
```

```
In [56]: conda install -c conda-forge imbalanced-learn
```

Collecting package metadata (current_repodata.json): done
Solving environment: done

==> WARNING: A newer version of conda exists. <==
current version: 22.11.1
latest version: 23.9.0

Please update conda by running

```
$ conda update -n base -c conda-forge conda
```

Or to minimize the number of packages updated during conda update use

```
conda install conda=23.9.0
```

Package Plan

environment location: /opt/conda

added / updated specs:
- imbalanced-learn

The following packages will be downloaded:

package	build		
-----	-----		
ca-certificates-2023.7.22	hbcca054_0	146 KB	conda-forge
certifi-2023.7.22	pyhd8ed1ab_0	150 KB	conda-forge
imbalanced-learn-0.11.0	pyhd8ed1ab_0	138 KB	conda-forge
openssl-3.1.4	hd590300_0	2.5 MB	conda-forge
-----	-----		
Total:		2.9 MB	

The following NEW packages will be INSTALLED:

```
imbalanced-learn  conda-forge/noarch::imbalanced-learn-0.11.0-pyhd8ed1ab_0
```

The following packages will be UPDATED:

```
ca-certificates          2022.12.7-ha878542_0 --> 2023.7.22-hbcca054_0
certifi                  2022.12.7-pyhd8ed1ab_0 --> 2023.7.22-pyhd8ed1ab_0
openssl                  3.0.7-h0b41bf4_1 --> 3.1.4-hd590300_0
```

Downloading and Extracting Packages

imbalanced-learn-0.1	138 KB		0%
ca-certificates-2023	146 KB		0%
certifi-2023.7.22	150 KB		0%

openssl-3.1.4	2.5 MB		0%
imbalanced-learn-0.1	138 KB	####2	12%
openssl-3.1.4	2.5 MB	2	1%
ca-certificates-2023	146 KB	####	11%
certifi-2023.7.22	150 KB	#####	100%
ca-certificates-2023	146 KB	#####	100%
imbalanced-learn-0.1	138 KB	#####	100%

Preparing transaction: done
 Verifying transaction: done
 Executing transaction: done

Note: you may need to restart the kernel to use updated packages.

```
In [60]: from imblearn import under_sampling, over_sampling
from imblearn.over_sampling import SMOTE
```

```
In [62]: X, y = X_trans[:, :-1], y
resampler = SMOTE(sampling_strategy="auto")
X_res, y_res = resampler.fit_resample(X, y)
```

```
In [63]: print("Before resampling, Shape of training instances: ", np.c_[X, y].shape)
print("After resampling, Shape of training instances: ", np.c_[X_res, y_res].shape)
```

Before resampling, Shape of training instances: (80, 465)
 After resampling, Shape of training instances: (148, 465)

```
In [64]: ## Target Cats after Resampling

print(np.unique(y_res))
print(f"Value Counts: \n-1: {len(y_res[y_res == -1])}, 1: {len(y_res[y_res == 1])}")

[-1  1]
Value Counts:
-1: 74, 1: 74
```

```
In [65]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=1/3, ra
```

```
print(f"train set: {X_train.shape, y_train.shape}")
print(f"test set: {X_test.shape, y_test.shape}")
```

```
train set: ((98, 464), (98,))
test set: ((50, 464), (50,))
```

In [66]: *# Prepared training and test sets*

```
X_prep = X_train
y_prep = y_train
X_test_prep = X_test
y_test_prep = y_test

print(X_prep.shape, y_prep.shape)
print(X_test_prep.shape, y_test_prep.shape)

(98, 464) (98,)
(50, 464) (50,)
```

In [68]: *!pip install xgboost*

```
Collecting xgboost
  Downloading xgboost-2.0.1-py3-none-manylinux2014_x86_64.whl (297.1 MB)
297.1 MB 6.0 MB/s eta 0:00:0000:0100:01
Requirement already satisfied: numpy in /opt/conda/lib/python3.10/site-packages (from xgboost) (1.23.5)
Requirement already satisfied: scipy in /opt/conda/lib/python3.10/site-packages (from xgboost) (1.9.3)
Installing collected packages: xgboost
Successfully installed xgboost-2.0.1
```

In [69]:

```
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import roc_auc_score
```

```
svc_clf = SVC(kernel='linear')
svc_rbf_clf = SVC(kernel='rbf')
random_clf = RandomForestClassifier(random_state=42)
xgb_clf = XGBClassifier()
```

In [70]: *## A function to display Scores*

```
def display_scores(scores):
    print("Scores: ", scores)
    print("Mean: ", scores.mean())
    print("Standard Deviation: ", scores.std())
```

In [71]: *## SVC Scores*

```
svc_scores = cross_val_score(svc_clf, X_prep, y_prep, scoring='roc_auc', cv=10, verbose=0)
display_scores(svc_scores)
```

```
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
Scores: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
Mean: 1.0
Standard Deviation: 0.0
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 0.1s finished
```

In [72]: *## Performance on test set using cross-validation*

```
# Predictions using cross-validation
svc_preds = cross_val_predict(svc_clf, X_test_prep, y_test_prep, cv=5)

# AUC score
svc_auc = roc_auc_score(y_test_prep, svc_preds)
svc_auc
```

Out[72]: 0.94

In [73]: *## SVC rbf Scores*

```
svc_rbf_scores = cross_val_score(svc_rbf_clf, X_prep, y_prep, scoring='roc_auc', cv
display_scores(svc_rbf_scores)
```

```
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
[CV] END ..... total time= 0.0s
Scores: [1. 1. 1. 1. 1. 0.96 1. 1. 1. 1. ]
Mean: 0.9960000000000001
Standard Deviation: 0.011999999999999976
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 0.1s finished
```

In [74]: *## Performance on test set using cross-validation*

```
# Predictions using cross-validation
svc_rbf_preds = cross_val_predict(svc_rbf_clf, X_test_prep, y_test_prep, cv=5)

# AUC score
```

```
svc_rbf_auc = roc_auc_score(y_test_prep, svc_rbf_preds)
svc_rbf_auc
```

Out[74]: 0.6799999999999999

In [77]: *## Random Forest Scores*

```
random_clf_scores = cross_val_score(random_clf, X_prep, y_prep, scoring='roc_auc',
display_scores(random_clf_scores))
```

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] END total time= 0.2s

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.2s remaining: 0.0s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

[CV] END total time= 0.2s

Scores: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

Mean: 1.0

Standard Deviation: 0.0

[Parallel(n_jobs=1)]: Done 10 out of 10 | elapsed: 2.3s finished

In [78]: *## Performance on test set using cross-validation*

Predictions using cross-validation

```
random_clf_preds = cross_val_predict(random_clf, X_test_prep, y_test_prep, cv=5)
```

AUC score

```
random_clf_auc = roc_auc_score(y_test_prep, random_clf_preds)
```

```
random_clf_auc
```

Out[78]: 1.0

In []: